

# Compose Chaos

Part 1: RAI Applied to Concurrency

<image of “chaos” in a lego piece shape>

## A. Joël Lamotte

- Experience:  
Games, Tools, Embedded, Mobile, Web...
- **C++**, Python, C#, JavaScript, Java, ActionScript...
- Today: **SoftBank Robotics Europe**  
(**libqi**, code “expert”, etc.)

**WE ARE RECRUITING!**

# Are you familiar with:

- mutex/locks
- atomics
- futures/promises, continuations, “.then()”
- “async(task)” -like functions
- concurrent containers
- executors
- thread-pools
- strands

# This talk is part of an ongoing exploration...

- ~~“expert friendly”~~
- ~~Parallelization~~
- “Heterogeneous” Concurrency
- Mostly pseudo code
- Note **slides numbers** for questions at the end

# “Heterogeneous” Concurrency

Different kinds of:

- concurrency
  - including “fork-join” parallelization
- data
- processing resources

Must **work together** in the same application

Most important rule when writing concurrent code:

**Do not share data.**

If you can.

You can't.

# Principles/Hypotheses

1. RAII applied to concurrency
2. No raw synchronization primitives
3. Task  $\Leftrightarrow$  message
4. System  $\Leftrightarrow$  “Actor-like”
5. 1 synchronization mechanism per abstraction
6. Executor( Task{} );

# Principles/Hypotheses

1. **RAII applied to concurrency**
2. No raw synchronization primitives
3. Task  $\Leftrightarrow$  message
4. System  $\Leftrightarrow$  “Actor-like”
5. 1 synchronization mechanism per abstraction
6. Executor( Task{} );



# RAII applied to concurrency

## Hypotheses:

- Constructor and destructor are **synchronization points**.
- Tasks launched by an object and accessing internal data of this object are **part of that object's lifetime**.

## RAII: Constructor and Destructor

1. Constructor: once constructed, an object must be able to handle concurrent member calls.
2. Destructor: all tasks launched by this object and accessing its internal data must be synchronized before the end of the destructor.

## RAII: Constructor and Destructor

1. Constructor: once constructed, an object must be able to handle concurrent member calls.
2. Destructor: all tasks launched by this object and accessing its internal data must be synchronized before the end of the destructor.

```
struct Data { int code; /* ... */ };

// might be concurrent
future<Data> crunch(int value);

struct System {
    explicit System(int x)
        : store{ crunch(x).get() }
    {}
    // ...
private:
    vector<Data> store;
    // ...
};
```

```
future<vector<int>> big_fat_random_garbage();

struct System {
    explicit System(int x) {
        auto ft_1 = crunch(some_cosmic_constant);
        auto ft_2 = crunch(x);
        auto ft_ints = big_fat_random_garbage();

        store = {ft_1.get(), ft_2.get()};
        values = ft_ints.get();
    }
    // ...
private:
    vector<Data> store;
    vector<int> values;
    // ...
};
```

```
int math_ninja_flip(const Data& x, const Data& y, int value);

struct System {
    explicit System(int x) {
        auto ft_1 = crunch(some_cosmic_constant);
        auto ft_2 = crunch(x);
        store = {ft_1.get(), ft_2.get()};

        values = { math_ninja_flip(store[0], store[1], x) };
    }
    // ...
private:
    vector<Data> store;
    vector<int> values;
    // ...
};
```

```
int math_ninja_flip(const Data& x, const Data& y, int value);

struct System {
    explicit System(int x)
        : store{crunch(some_cosmic_constant).get(), crunch(x).get()}
        , values{ math_ninja_flip(store[0], store[1], x) }
    {
    }
    // ...
private:
    vector<Data> store;
    vector<int> values;
    // ...
};
```

## RAII: Constructor

After the end of the constructor, the object must be able to handle concurrent member calls.

- Initialization must be synchronous, even if blocking is necessary (!!)
  - Don't Block™?



```
struct System {  
    explicit System(int x) {  
        auto ft_data2 = crunch(some_cosmic_constant);  
        crunch(x).then(=[](auto ft_data){  
            store.push_back( ft_data.get() );  
        });  
  
        big_fat_random_garbage().then([&](auto ft_ints){  
            values = ft_ints.get();  
        });  
    }  
    // ...  
private:  
    vector<Data> store;  
    vector<int> values;  
    // ...  
};
```

## RAII: Constructor

After the end of the constructor, the object must be able to handle concurrent member calls.

- Initialization must be synchronous
- Do not defer initialization after constructor's end

## RAII: Constructor

After the end of the constructor, the object must be able to handle concurrent member calls.

- Initialization must be synchronous
- Do not defer initialization after constructor's end?

```
struct System {
    explicit System(int x) {
        auto ft_1 = crunch(some_cosmic_constant); // Takes a long time...
        auto ft_2 = crunch(x); // Takes a long time...
        auto ft_ints = big_fat_random_garbage(); // Takes a long time...

        store = {ft_1.get(), ft_2.get()}; // Wait for a long time...
        values = ft_ints.get(); // Might wait a long time...
    }

private:
    vector<Data> store;
    vector<int> values;
    // ...
};
```

# RAII: Constructor

After the end of the constructor, the object must be able to handle concurrent member calls.

- Initialization must be synchronous
- ~~Do not defer initialization after constructor's end...?~~
- Exit the constructor as soon as ready to handle concurrent member calls?

```

struct System {
    explicit System(int x) {
        crunch(x).then([&]( auto ft_data ){ // *
            scoped_lock lock{this_mutex};
            store.push_back( ft_data.get() );
        });
        big_fat_random_garbage()
            .then([&](auto ft_ints){ // *
                scoped_lock lock{this_mutex};
                values.push_back( ft_ints.get() );
            });
    }
    // ...
private:
    vector<Data> store;
    vector<int> values;
    mutex this_mutex;
    // ...
};

```

```

future<int> System::foowizz(Data data) {
    return async( [this, data]{ // *
        scoped_lock lock{this_mutex};
        int result = math_ninja_flip(
            store.back(), data, values.back());
        store.push_back(data);
        values.push_back(data.code);
        return result;
    });
}

// * : Tasks are not guaranteed to be executed
// at a time where the object is still alive!
// Also initialization may finish after
// other tasks are executed!

```

```
System sys_a{1234};  
sys_a.foowizz(some_data);
```

```
System sys_b{23423};  
sys_b.foowizz(some_data);
```

```
System sys_c{4329423};  
sys_c.foowizz(some_data);
```

```
System sys_a{1234};
System sys_b{23423};
System sys_c{4329423};
// "Simplified" version
auto ft_data =
    sys_a.foowizz(some_data)
    .then([](int value){
        return crunch(value);
    })
    .then([&](Data data){
        return sys_b.foowizz(data);
    })
    .then([](int value){
        return crunch(value);
    })
    .then([&](Data data){
        return sys_c.foowizz(data);
    });
```



# RAII: Constructor

After the end of the constructor, the object must be able to handle concurrent member calls.

- Initialization must be synchronous
- Do not defer initialization after constructor's end
- **Optimization**: Exit the constructor as soon as ready to handle concurrent member calls

```

struct System {
    explicit System(int x) {
        auto ft_data = crunch(x);
        auto ft_ints = big_fat_random_garbage();
        work_queue.push( [=]{
            store.push_back(ft_data.get());
            values = ft_ints.get();
        });
    }
    future<int> foowizz(Data data); // ...
private:
    vector<Data> store;
    vector<int> values;
    concurrent_queue<function<void()>>
        work_queue;
    thread update_thread{ update_loop };
    void update_loop(); // pop and execute tasks
};

```

```

future<int> System::foowizz(Data data) {
    return work_queue.push( [this, data]{
        int result = math_ninja_flip(
            store.back(), data,
            values.back());
        store.push_back(data);
        values.push_back(data.code);
        return result;
    });
}

```

```

struct System {
    explicit System(int x) {
        auto ft_data = crunch(x);
        auto ft_ints = big_fat_random_garbage();
        strand.push( [=]{
            store.push_back(ft_data.get());
            values = ft_ints.get();
        });
    }
    future<int> foowizz(Data data); // ...
private:
    vector<Data> store;
    vector<int> values;
    Strand strand;
    // ...
};

```

```

future<int> System::foowizz(Data data) {
    return strand.push( [this, data]{
        int result = math_ninja_flip(
store.back(), data, values.back());
        store.push_back(data);
        values.push_back(data.code);
        return result;
    });
}

```

# RAII: Constructor and Destructor

1. Constructor: once constructed, object must be able to handle concurrent member calls.
2. Destructor: all tasks launched by this object and accessing its internal data must be synchronized before the end of the destructor.

# RAII: Destructor

Before the end of the destructor, all tasks launched accessing internal data must **either**:

- **Be finished**
- **Do nothing on execution**
  - Don't throw!

## RAII: Destructor

- Wait for tasks to finish

```

class Poppy {

    template<class F>
    void launch(F task) {
        launched_tasks.push(async(executor, task));
    }

public:
    Poppy(Executor exec) : executor(exec) {
        launch([]{ init_background_system(); });
    }

    ~Poppy() {
        future<void> ft;
        while(launched_task.pop(ft))
            ft.wait();
    }
}

```

```

int count() const { return counter; }

void im_poppy(){
    launch([&]{
        some_dirty_laundry();
        ++counter;
    });

private:
    atomic<int> counter;
    Executor executor;
    concurrent_queue<future<void>>
        launched_tasks;
};

```

```
class Poppy {  
    //...  
    ~Poppy() {  
        future<void> ft;  
        while(launched_task.pop(ft))  
            ft.wait();  
    } // ...  
};
```

```
class BigSystem {  
    ThreadPool thread_pool;  
    Poppy poppy{ &thread_pool };  
    // ...  
public:  
    ~BigSystem() {  
        thread_pool.stop_tasks_execution();  
        some_cleanup();  
    } // ~Poppy() might be blocked!  
  
    void be_funny() { poppy.im_poppy(); }  
    // ...  
};
```



```
class Poppy {  
    //...  
    ~Poppy() {  
        future<void> ft;  
        while(  
            launched_task.pop(ft)  
        ) ft.wait();  
    } // ...  
};
```

```
void be_poppy_for_some_time(const seconds timeout){  
    concurrent_queue<function<void()>> work_queue;  
    Poppy poppy{ [&](auto task){  
        work_queue.push(task);  
    }};  
  
    auto begin_time = some_clock::now();  
    begin_try_to_be_poppy(poppy); // calls poppy's members a lot  
    while((some_clock::now() - begin_time) < timeout) {  
        function<void()> task;  
        if(work_queue.pop(task))  
            task();  
    } // No more task execution  
    end_try_to_be_poppy(poppy);  
} // ~Poppy() might be blocked!
```

```
class Poppy {  
    //...  
    ~Poppy() {  
        future<void> ft;  
        while(  
            launched_task.pop(ft)  
        ) ft.wait();  
    } // ...  
};
```

```
void be_poppy_for_some_time(const seconds timeout){  
    Strand strand;  
    Poppy poppy{ &strand };  
  
    begin_try_to_be_poppy(poppy); // calls poppy's members a lot  
    sleep_for(timeout);  
    strand.join(); // stop executing tasks, MAYBE destroy tasks?  
    end_try_to_be_poppy(poppy);  
} // ~Poppy() might be blocked!  
    // Depends on what the executor guarantees...
```

## RAII: Destructor

- ~~Wait for tasks to finish~~
- Notify tasks to do nothing on execution

```

class Poppy {
    template<class F> void launch(F f) {
        async(executor, [=]{
            if(tasks_must_not_execute) // might not exist!
                f();
        });
    }
public:
    ~Poppy() {
        tasks_must_not_execute = true;
        // don't wait for running tasks...
    }
    void im_poppy() {
        launch([this]{
            some_dirty_laundry();
            ++counter; // might not exist!
        });
    }
}

```

```

// ...
private:
    atomic<int> counter{0};
    atomic<bool>
        tasks_must_not_execute{false};
    // ...
};

```

```

class Poppy {
    template<class F> void launch(F f) {
        weak_ptr maybe_state{state}; // not thread-safe!
        async(executor, [=]{
            if(auto state = maybe_state.lock())
                f(state);
        });
    }
public:
    ~System() = default; // no wait

    void im_poppy(){
        launch([](auto state){
            some_dirty_laundry();
            ++state->counter; // forced "living" after destruction
        });
    }
    // ...

```

```

private:
    struct State {
        atomic<int> counter{0};
    };

    // needs protection too!
    shared_ptr<State> state =
        make_shared<State>();

    // ...
};

```

## RAII: Destructor

- ~~Wait for tasks to finish~~
- ~~Notify tasks to do nothing on execution~~
- Wait running tasks to finish?

## RAII: Destructor - The Right Way\*

1. Notify all tasks to “join”;

- If a task did not execute yet, make it do nothing on execution;

2. Wait already running tasks to finish;

- Stop waiting as soon as you can ignore tasks' end;

\* Except in critical cases where you have to execute all tasks anyway: banking, critical logging, black boxes systems for exemple.

## RAII: Destructor

1. Notify **all tasks** to “join”;
  - If a task did not execute yet, make it do nothing on execution;
2. Wait already running tasks to finish;
  - Stop waiting as soon as you can ignore tasks end;

**Keep track of internal tasks' lifetime**



## RAII: Destructor

1. Notify all tasks to “join”;
  - If a task did not execute yet, make it do nothing on execution;
2. Wait already **running tasks** to finish;
  - Stop waiting as soon as you can ignore tasks end;

**Keep track of internal tasks' status**

## RAII: Destructor

1. **Notify** all tasks to “join”;
  - If a task did not execute yet, make it do nothing on execution;
2. **Wait** already running tasks to finish;
  - Stop waiting as soon as you can ignore tasks end;

**Use a higher-level task synchronization tool  
(could be a strand or something similar...)**

```

class Poppy {
    template<class F> void launch(F task) {
        auto synchronized_task = task_sync.synced(task);
        async(executor, synchronized_task);
    }
public:
    ~System() = default;
    // task_sync.join() called in ~TaskSynchronizer()
    // wait for running tasks, deactivate all other tasks

    void im_poppy(){
        launch([&]{
            some_dirty_laundry();
            ++counter; // 'this' is guaranteed to be alive
                      // if this code is executing
        });
    }
    // ...

```

```

private:
    atomic<int> counter;
    // ...
    TaskSynchronizer task_sync;
};

```

## RAII: Destructor

1. Notify all tasks to “join”;
  - If a task did not execute yet, make it do nothing on execution;
2. Wait already running tasks to finish;
  - Stop waiting as soon as you can ignore tasks end;
3. Do not wait asynchronous calls from the destructor: fire and forget?

## RAII: Destructor

1. Notify all tasks to “join”;
  - If a task did not execute yet, make it do nothing on execution;
2. Wait already running tasks to finish;
  - Stop waiting as soon as you can ignore tasks end;
3. Do not wait asynchronous calls from the destructor: ~~fire and forget~~. If reasonable?

# RAII: Destructor

Do not wait asynchronous calls from the destructor: fire and forget, if reasonable.

```
struct System {  
    // ...  
    World& world; // dependency, not owned  
public:  
    System(World& w) : world(w) {}  
    ~System()  
    {  
        world.notify_everybody(SystemEnd{}).get();  
    }  
    // ...  
};
```

```
struct World {  
    // ...  
    template<class M>  
    future<void> notify_everybody(M message);  
};
```

# RAII: Destructor

Do not wait asynchronous calls from the destructor: fire and forget, If reasonable.

```
struct System {  
    // ...  
    World& world; // dependency, not owned  
public:  
    System(World& w) : world(w) {}  
    ~System()  
    {  
        world.notify_everybody(SystemEnd{});  
        // fire and forget: ok  
    }  
    // ...  
};
```

```
struct World {  
    // ...  
    template<class M>  
    future<void> notify_everybody(M message);  
};
```

# RAII: Destructor

Do not wait asynchronous calls from the destructor: fire and forget, if reasonable.

```
class Focus { // Unique focus acquisition
    SomeSystem& system;
public:
    // ... move-only
    Focus(SomeSystem& sys)
        : system{ sys }
    { system.acquire_focus().get(); } // wait to get it
    // fire and forget: might be ok, we don't really own or share the resource
    ~Focus() { system.notify_focus_end_use(); } // async call, no wait
    // ...
};
```



# RAII: Constructor and Destructor

Observations:

```
template<class T> class Mechanism { /* ...  
*/ };
```

```
class SubSystem {  
    BigData data;  
    Mechanism cog{data};  
    // ...  
};
```

```
class BigSystem {  
    ThreadPool executor;  
    SubSystem a{executor};  
    OtherSubSystem b{a, executor};  
    WeirdSubSystem c{b, executor};  
    // ...  
};
```

```
class VeryBigSystem  
{  
    BigSystem k;  
    OtherBigSystem l;  
    TroublingSystem m{l, k};  
    // ...  
};
```

```

class Client
{
public:
    Client( std::string window_name, Configuration
configuration );
    ~Client() = default;

    core::TaskChainThread io_executor;

    core::SystemIO system_io{ io_executor };
    core::WindowSystem window_io{ system_io };
    core::Window main_window;

    input::InputEngine input_engine;
    graphic::GraphicEngine graphic_engine;
    GraphicComponents graphic_components;
    ActionsRegistry actions_registry;
    // ...
};

```

```

Client::Client( std::string window_name,
Configuration config )
    : io_executor( IO_CYCLE_FREQUENCY )
    , main_window( window_io, window_name,
        core::read_window_state( config ) )
    , input_engine{ system_io, main_window,
        input_config( config ) }
    , graphic_engine{ std::ref(io_executor),
        main_window, config }
    , graphic_components{ graphic_engine }
    , actions_registry{ input_engine }
{
    main_window.show();
}

```

# RAII: Constructor and Destructor

## Observations:

- Like for all RAII-based code, reasoning concurrent code appears “easy”, familiar, unsurprising.
- Helps with correctness.
- Opens door to potential optimizations (by avoiding blocking internally once algorithms are correct).

## Principles/Hypotheses

1. RAII applied to concurrency
2. No raw synchronization primitives
3. Task  $\Leftrightarrow$  message
4. System  $\Leftrightarrow$  “Actor-like”
5. 1 synchronization mechanism per abstraction
6. Executor( Task{} );

# RAII applied to Concurrency

## Questions?

### References:

Sean Parent's concurrency talk: <https://www.youtube.com/watch?v=32f6JrQPv8c>

Book "C++ Concurrency In Action":

<https://www.manning.com/books/c-plus-plus-concurrency-in-action>

@mjklaim [mjklaim@gmail.com](mailto:mjklaim@gmail.com) or [jlamotte@softbankrobotics.com](mailto:jlamotte@softbankrobotics.com)

# RAII applied to Concurrency

BONUS CODE!

TaskSynchronizer

# Projects & Experimentations

## My Technical Goals:

1. code maintainable for **long-term**
2. **scalable** usage of available resources
3. good **performance**
4. open door to easy **extensibility**



# Projets & Experimentations

**Qi library** (SBE): <https://github.com/aldebaran/libqi>

An inter-process RPC-based communication middleware and concurrency coding tools.

- Targets robotics use cases (multi-process)
- Gives good perspective vs experimentations

**WE ARE STILL RECRUITING**